

HW 4

SDS348 Spring 2021

Name: Santhosh Saravanan

EID: sks3648

This homework is due on Mar 8, 2021 at 8am. Submit a pdf file on Gradescope.

For all questions, include the R commands/functions that you used to find your answer (show R chunk).

Answers without supporting code will not receive credit. Write full sentences to describe your findings.

Question 1: (9 pts)

The dataset `world_bank_pop` is a built-in dataset in `tidyverse`. It contains information about total population and population growth, overall and more specifically in urban areas, for countries around the world.

1.1 (1 pt) Save the dataset `world_bank_pop` as `myworld` and take a look at it with `head()`. Is the data tidy? Why or why not?

```
library(tidyverse)
myworld <- world_bank_pop
head(myworld)
```

```
## # A tibble: 6 x 20
##   country indicator `2000` `2001` `2002` `2003` `2004` `2005` `2006`
##   <chr>    <chr>      <dbl> <dbl> <dbl> <dbl> <dbl> <dbl> <dbl>
## 1 ABW     SP.URB.TO... 4.24e4 4.30e4 4.37e4 4.42e4 4.47e+4 4.49e+4 4.49e+4
## 2 ABW     SP.URB.GR... 1.18e0 1.41e0 1.43e0 1.31e0 9.51e-1 4.91e-1 -1.78e-2
## 3 ABW     SP.POP.TO... 9.09e4 9.29e4 9.50e4 9.70e4 9.87e+4 1.00e+5 1.01e+5
## 4 ABW     SP.POP.GR... 2.06e0 2.23e0 2.23e0 2.11e0 1.76e+0 1.30e+0 7.98e-1
## 5 AFG     SP.URB.TO... 4.44e6 4.65e6 4.89e6 5.16e6 5.43e+6 5.69e+6 5.93e+6
## 6 AFG     SP.URB.GR... 3.91e0 4.66e0 5.13e0 5.23e0 5.12e+0 4.77e+0 4.12e+0
## # ... with 11 more variables: 2007 <dbl>, 2008 <dbl>, 2009 <dbl>, 2010 <dbl>,
## #   2011 <dbl>, 2012 <dbl>, 2013 <dbl>, 2014 <dbl>, 2015 <dbl>, 2016 <dbl>,
## #   2017 <dbl>
```

Each country abbreviation forms a row with information stemming from the indicator and year. Each variable (indicator, year, and country code) all form a distinct column. Hence, the data is tidy.

1.2 (1 pt) Using pipes and `dplyr` functions, how many countries are there in the dataset?

```
# your code goes here (make sure to add comments)
myworld %>%
  summarise(count=n_distinct(country))
```

```
## # A tibble: 1 x 1
##   count
##   <int>
## 1    264
```

#first we group by country and use count parameter in summarise dplyr method

There are 264 countries in the dataset.

1.3 (2 pts) Use one of the `pivot` functions to create a new dataset, `myworld2`, with the years 2000 to 2017 appearing as a **numeric** variable **year**, and the different values for the indicator variable are in a variable called **value**. In this new dataset, how many lines are there per country? Why does it make sense?

```
# your code goes here (make sure to add comments)
myworld2 <- myworld %>% pivot_longer(-c(country,indicator), names_to = "year", names_
transform = list(year= as.integer))
head(myworld2)
```

```
## # A tibble: 6 x 4
##   country indicator    year value
##   <chr>    <chr>    <int> <dbl>
## 1 ABW      SP.URB.TOTL  2000  42444
## 2 ABW      SP.URB.TOTL  2001  43048
## 3 ABW      SP.URB.TOTL  2002  43670
## 4 ABW      SP.URB.TOTL  2003  44246
## 5 ABW      SP.URB.TOTL  2004  44669
## 6 ABW      SP.URB.TOTL  2005  44889
```

```
myworld2%>%count(country,indicator,sort = TRUE) # find the number of times a country
is used for every type of indicator
```

```
## # A tibble: 1,056 x 3
##   country indicator      n
##   <chr>    <chr>    <int>
## 1 ABW      SP.POP.GROW    18
## 2 ABW      SP.POP.TOTL    18
## 3 ABW      SP.URB.GROW    18
## 4 ABW      SP.URB.TOTL    18
## 5 AFG      SP.POP.GROW    18
## 6 AFG      SP.POP.TOTL    18
## 7 AFG      SP.URB.GROW    18
## 8 AFG      SP.URB.TOTL    18
## 9 AGO      SP.POP.GROW    18
## 10 AGO     SP.POP.TOTL    18
## # ... with 1,046 more rows
```

```
myworld2%>%count(country,sort = TRUE) # just to confirm that the total number of rows
```

```
## # A tibble: 264 x 2
##   country      n
##   <chr>    <int>
## 1 ABW        72
## 2 AFG        72
## 3 AGO        72
## 4 ALB        72
## 5 AND        72
## 6 ARB        72
## 7 ARE        72
## 8 ARG        72
## 9 ARM        72
## 10 ASM        72
## # ... with 254 more rows
```

```
# for every country is 72.
```


There are about 72 lines per country. This makes sense because there are 4 indicators: population growth, total population, urban population growth, and total urban population. Given that the years range from 2000 to 2017, that's about 17 years. 18 years per indicator gives about 72 measurements per country, so it makes sense.

1.4 (3 pts) Represent the total population growth and urban population growth in the world (country code is WLD) between 2000 and 2017. How has population growth changed over the years?

```
# your code goes here (make sure to add comments)
myworld2%>% filter(country=="WLD" & indicator == "SP.URB.GROW")
```

```
## # A tibble: 18 x 4
##   country indicator   year value
##   <chr>    <chr>      <int> <dbl>
## 1 WLD      SP.URB.GROW  2000  2.16
## 2 WLD      SP.URB.GROW  2001  2.27
## 3 WLD      SP.URB.GROW  2002  2.35
## 4 WLD      SP.URB.GROW  2003  2.33
## 5 WLD      SP.URB.GROW  2004  2.32
## 6 WLD      SP.URB.GROW  2005  2.31
## 7 WLD      SP.URB.GROW  2006  2.28
## 8 WLD      SP.URB.GROW  2007  2.24
## 9 WLD      SP.URB.GROW  2008  2.26
## 10 WLD     SP.URB.GROW  2009  2.23
## 11 WLD     SP.URB.GROW  2010  2.20
## 12 WLD     SP.URB.GROW  2011  2.08
## 13 WLD     SP.URB.GROW  2012  2.08
## 14 WLD     SP.URB.GROW  2013  2.07
## 15 WLD     SP.URB.GROW  2014  2.06
## 16 WLD     SP.URB.GROW  2015  2.05
## 17 WLD     SP.URB.GROW  2016  2.03
## 18 WLD     SP.URB.GROW  2017  2.00
```


```
ggplot(myworld2%>% filter(country=="WLD" & indicator == "SP.URB.GROW"),aes(x= `year`,
y= `value`)) + geom_point(size=1.5,color="Blue") + geom_line() + labs(y="Urban popula
tion growth",x="Year") + ggtitle("Urban Population Growth by Year in World")
```



Year

```
ggplot(myworld2%>% filter(country=="WLD" & indicator == "SP.POP.GROW"),aes(x= `year`,
y= `value`)) + geom_point(size=1.5,color="Blue") + geom_line() + labs(y="Total popula
tion growth",x="Year") + ggtitle("Total Population Growth by Year in World")
```

Total Population Growth by Year in World



Urban population growth reached its peak above 2.3 at around 2002-2003 and has declined ever since until past 2005. After 2005, there was a momentary spike upward at around 2007-2007 but has declined until experiencing a sharp decline after 2010. Ever since, the rate of decrease is still present, but a lot slower after 2011. For total population growth, it's been steadily decreasing till a a short while after 2010. There was an all-time low at the time at 2011, but it steadily gre back and has similarly, the rate of decrease is still present, but it's a bit higher than the urban population growth after 2011.

1.5 (2 pts) Use one of the `pivot` functions to create a new dataset, `myworld3`, with the different categories for the indicator variable appearing as their own variables. Use `dplyr` functions to rename `SP.POP.GROW` and `SP.URB.GROW`, as `pop_growth` and `pop_urb_growth` respectively. What is the country code that had the highest population growth in 2017?

```
myworld3 <- myworld2 %>% pivot_wider( names_from= indicator, values_from= value)
myworld3 <- myworld3 %>% rename(pop_growth = SP.POP.GROW, pop_urb_growth = SP.URB.GRO
W)
myworld3 %>% filter(year==2017) %>% arrange(desc(pop_growth))
```

```
## # A tibble: 264 x 6
##   country year SP.URB.TOTL pop_urb_growth SP.POP.TOTL pop_growth
##   <chr>   <int>      <dbl>          <dbl>      <dbl>      <dbl>
## 1 OMN     2017    3874061          5.95    4636262      4.67
## 2 BHR     2017    1331176          4.73    1492584      4.62
## 3 NRU     2017     13649          4.50     13649      4.50
## 4 NER     2017    3511546          4.18    21477348      3.82
## 5 GNQ     2017     908248          4.42    1267689      3.71
## 6 AGO     2017   19311773          4.38   29784193      3.31
## 7 UGA     2017    9942492          5.76   42862958      3.26
## 8 COD     2017   35691987          4.57   81339988      3.25
## 9 BDI     2017   1380411          5.72   10864245      3.18
## 10 TZA    2017   18942681          5.28   57310019      3.08
## # ... with 254 more rows
```

OMN has the highest population growth in 2017.

Question 2: (10 pts)

From answering the previous question, we have no idea what actual countries are represented by the codes. We will now use a package that has information about the coding system used by the World bank.

2.1 (2 pts) Install the package `countrycode`. We will use a built-in dataset called `codelist`. Make sure to upload the library and save this dataset as `mycodes`. Using `dplyr` functions, modify `mycodes` to: 1. select only the variables `continent`, `wb` (World Bank code), and `country.name.en` (country name in English); 2. filter to keep countries in Europe only; 3. remove countries with missing `wb` code. How many countries are there in Europe with a World Bank code?

```
#install.packages("countrycode")

library(countrycode)

mycodes <- codelist
mycodes <- mycodes %>% select(continent, wb, country.name.en) %>% filter(continent ==
"Europe") %>% drop_na(wb) # select countries based on continent, wb, and country name
in English
# filter with continent == "Europe" and drop_na values in wb column
nrow(mycodes)
```

```
## [1] 46
```

```
# your code goes here (make sure to add comments)
```

There are 46 countries in Europe with a World Bank Code.

2.2 (2 pts) Use a `left_join()` function to create a new dataset, `myeurope`, to add data to the countries in `mycodes` dataset from `myworld3` dataset. Match the two datasets based on the World Bank code. Using `dplyr` functions, change the name of the variable containing the World

Bank code to country .

```
# your code goes here (make sure to add comments)
mycodes <- rename(mycodes,c("country" = "wb"))
myeurope <- mycodes %>%
left_join(myworld3, by= c("country"))
```

2.3 (1 pt) Using `dplyr` functions, what was the total population in European countries in 2017? Give your answer in million (round to the next million).

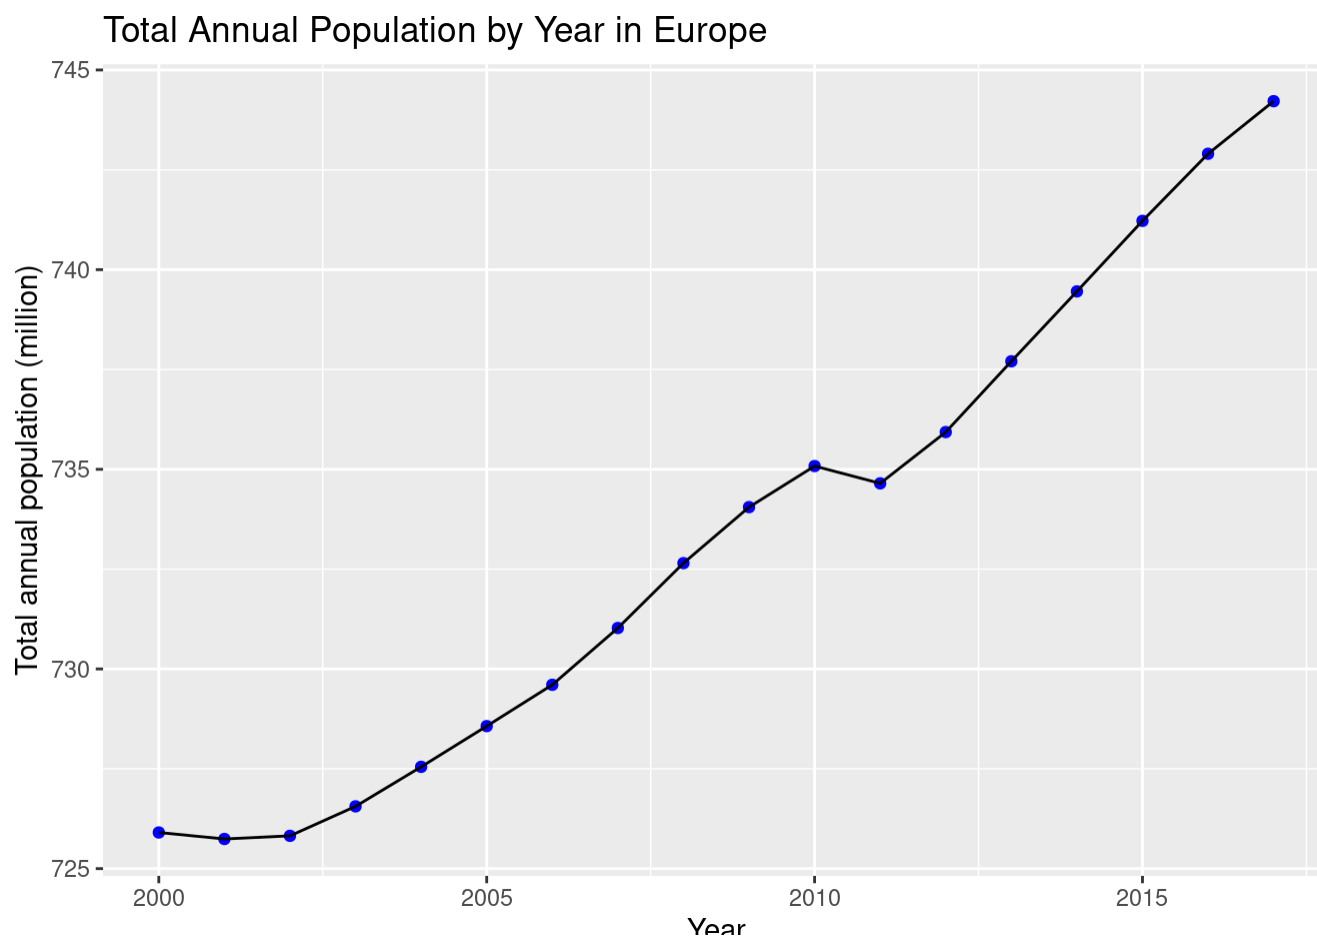
```
myeurope1 <- myeurope %>% filter(continent == 'Europe' & year == 2017)
sum(myeurope1$SP.POP.TOTL)
```

```
## [1] 744218594
```

There are 744 million people in Europe in 2017.

2.4 (2 pts) Represent the annual total population in European countries between 2000 and 2017. Express the total population in million. How has European population changed over the years?

```
annualNumbers = myeurope %>% group_by(year) %>% summarise(totalPop=sum(SP.POP.TOTL)/1
000000)
ggplot(annualNumbers,aes(x= `year`, y=totalPop )) + geom_point(size=1.5,color="Blue")
+ geom_line() + labs(y="Total annual population (million)",x="Year") + ggtitle("Total
Annual Population by Year in Europe")
```



Population seems to be strictly increasing except for 2011, where there was a slight dip. After 2011, it seems to continue to increase.

2.5 (2 pts) Create a new dataset `myeurope2017` by filtering the data for the year 2017, dropping the variable `year`, and creating a new variable `prop_urb` which is the proportion of urban population for each country. Which European country had the lowest proportion of urban population in 2017?

```
myeurope2017 <- myeurope %>%filter(year == 2017)%>% select(-year) %>% mutate(prop_urb
= SP.URB.TOTL/SP.POP.TOTL) %>% arrange(prop_urb)
myeurope2017
```

```
## # A tibble: 46 x 8
##   continent country country.name.en      SP.URB.TOTL pop_urb_growth SP.POP.TOTL
##   <chr>      <chr>      <chr>          <dbl>          <dbl>          <dbl>
## 1 Europe    LIE      Liechtenstein      5429          0.777      37922
## 2 Europe    FRO      Faroe Islands      20659          0.675      49290
## 3 Europe    MDA      Moldova            1510667        0.0461     3549750
## 4 Europe    BIH      Bosnia & Herzegovina 1679019        0.472     3507017
## 5 Europe    IMN      Isle of Man         44215          0.875       84287
## 6 Europe    SVK      Slovakia            2923996        0.0650     5439892
## 7 Europe    ROU      Romania            10564196       -0.523     19586539
## 8 Europe    SVN      Slovenia            1121686          0.550     2066748
## 9 Europe    SRB      Serbia              3928397       -0.276     7022268
## 10 Europe   HRV      Croatia             2337910       -0.705     4125700
## # ... with 36 more rows, and 2 more variables: pop_growth <dbl>, prop_urb <dbl>
```

Liechtenstein had the lowest proportion of urban population in 2017

2.6 (1 pt) Using `dplyr` functions, find the top 3 countries in terms of their total population in 2017.

```
topmyeurope <- myeurope %>%filter(year==2017)
arrange(topmyeurope,desc(SP.POP.TOTL))
```



```
## # A tibble: 46 x 8
##   continent country country.name.en year SP.URB.TOTL pop_urb_growth
##   <chr>      <chr>      <chr>      <int>      <dbl>      <dbl>
## 1 Europe    RUS      Russia      2017    107348258      0.278
## 2 Europe    DEU      Germany      2017     63890984      0.468
## 3 Europe    FRA      France      2017     53815732      0.715
## 4 Europe    GBR      United Kingdom 2017     54892898      0.958
## 5 Europe    ITA      Italy      2017     42473185      0.287
## 6 Europe    ESP      Spain      2017     37294880      0.489
## 7 Europe    UKR      Ukraine      2017     31043784     -0.253
## 8 Europe    POL      Poland      2017     22825379     -0.106
## 9 Europe    ROU      Romania      2017     10564196     -0.523
## 10 Europe   NLD      Netherlands 2017     15604089      1.09
## # ... with 36 more rows, and 2 more variables: SP.POP.TOTL <dbl>,
## #   pop_growth <dbl>
```

Russia, Germany, and France are the top 3 countries in terms of total population in 2017

Question 3: (6 pts)

When dealing with location data, we can actually visualize information on a map if we have geographic information such as latitude and longitude.

3.1 (1 pt) We will use a built-in function called `map_data()` to get geographic coordinates about countries in the world (see below). Take a look at the dataset `mapWorld` with `glimpse()`. What variable could we use to join this dataset with `myeurope2017` dataset?

```
# geographic coordinates about countries in the world
mapWorld <- map_data("world")
glimpse(mapWorld)
```

```
## Rows: 99,338
## Columns: 6
## $ long      <dbl> -69.89912, -69.89571, -69.94219, -70.00415, -70.06612, -70.0...
## $ lat       <dbl> 12.45200, 12.42300, 12.43853, 12.50049, 12.54697, 12.59707, ...
## $ group     <dbl> 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 2, 2, 2, 2, 2, 2, 2, 2, 2, ...
## $ order     <int> 1, 2, 3, 4, 5, 6, 7, 8, 9, 10, 12, 13, 14, 15, 16, 17, 18, 1...
## $ region    <chr> "Aruba", "Aruba", "Aruba", "Aruba", "Aruba", "Aruba", "Aruba...
## $ subregion <chr> NA, NA, NA, NA, NA, NA, NA, NA, NA, NA, NA, NA, NA, NA, NA, ...
```

We could use the region variable to join the dataset with the `myeurope2017` dataset.

3.2 (1 pt) We want to use a `left_join()` function to create a new dataset, `mymap`, to add data to the countries in `myeurope2017` dataset from `mapWorld` dataset, matching the two datasets based on the country name. If we then use `dplyr` functions, we can identify some missing values for `lat` and `long` in the new dataset. Indeed, some countries such as United Kingdom did not have a match. Why do you think this happened?

```
myeurope2017_1 <- rename(myeurope2017,c("region" = "country.name.en"))
mymap <- myeurope2017_1 %>%
left_join(mapWorld, by= c("region"))
```

For the `mapWorld` dataset, the country United Kingdom is saved as UK under the region. For the `myeurope2017` dataset, the country United Kingdom is saved as United Kingdom. When you left-join these two datasets, the names don't match.

3.3 (1 pt) To identify all countries that did not have an exact match, do an `anti_join()` and display only distinct country names. How many countries did not have an exact match? *Note: using `anti_join()` is a very useful function to identify differences between datasets.*

```
anti_join(myeurope2017_1,mapWorld,by="region")
```

```
## # A tibble: 5 x 8
##   continent country region      SP.URB.TOTL pop_urb_growth SP.POP.TOTL pop_growth
##   <chr>      <chr>   <chr>          <dbl>          <dbl>          <dbl>      <dbl>
## 1 Europe    BIH      Bosnia & ...  1679019         0.472        3507017    -0.279
## 2 Europe    MKD      North Mac...  1202983         0.415        2083160     0.0938
## 3 Europe    CZE      Czechia       7803157         0.379       10591323     0.236
## 4 Europe    GBR      United Ki...  54892898        0.958       66022273     0.648
## 5 Europe    GIB      Gibraltar     34571          0.473        34571       0.473
## # ... with 1 more variable: prop_urb <dbl>
```

5 countries did not have a match.

3.4 (1 pt) Joining datasets by variables containing names often leads to a mismatch because spelling can vary from one dataset to another. Sometimes we need to manually fix spelling in order to be able to match values. Consider the code given below. Replace the name of United Kingdom so that its name in `myeurope2017` dataset corresponds to the name given in `mapWorld` dataset. Following this code, add a pipe and use a `left_join()` function to create the new dataset, `mymap`, adding data to the countries in `myeurope` dataset from `mapWorld` dataset.

```
# myeurope2017_1 <- rename(myeurope2017,c("region" = "country.name.en"))
# mymap <- myeurope2017_1 %>%
# left_join(mapWorld, by= c("region"))
mymap <- myeurope2017 %>%
  mutate(country_clean = recode(country.name.en,
                                'United Kingdom' = 'UK',
                                'Bosnia & Herzegovina' = 'Bosnia and Herzegovina',
                                'Czechia' = 'Czech Republic',
                                'North Macedonia' = 'Macedonia'))%>% left_join(mapWorld,by= c("country.name.en" = "region"))

#
```

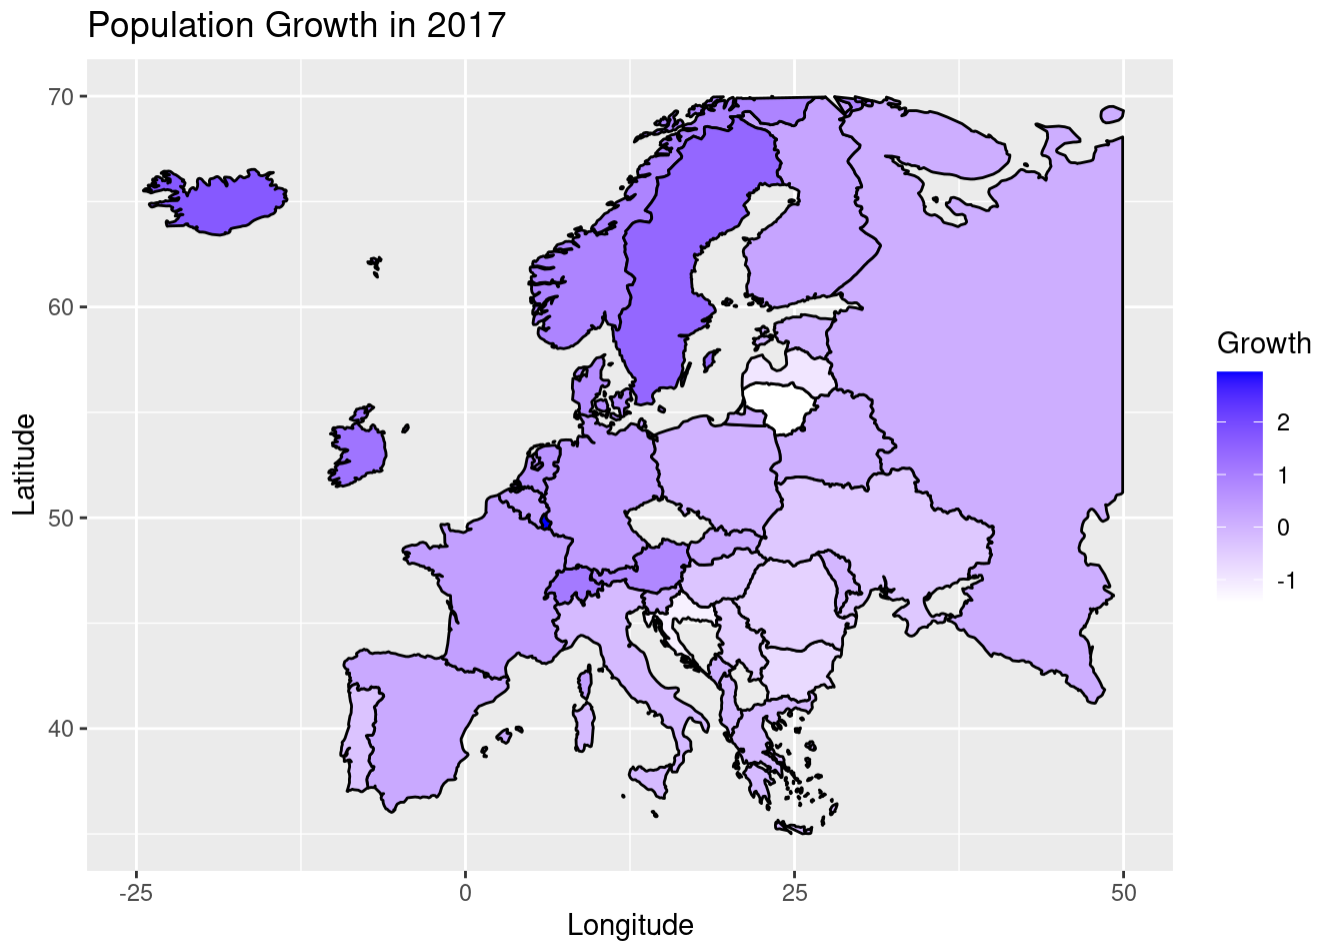
3.5 (2 pts) Let's visualize how population growth varies across European countries. Install the package `ggmap`, call the corresponding library, and use the R code provided below. Try to identify what each component of the graph does by completing the code with comments.

```

#install.packages("ggmap")
library(ggmap)

mymap %>%
  ggplot(aes(x=long, y=lat, group = group, fill = pop_growth)) +
  # Plots each individual country based on latitude and longitude in the ggplot and
  # outlines each country with a black color
  geom_polygon(colour = "black") +
  # Creates a smooth gradient between white and blue and colorcodes the bar relating
  # to # pop_growth. Higher pop_growth has colors close to blue. Lower pop_growth has
  # colors # close to white.
  scale_fill_gradient(low = "white", high = "blue", guide="colorbar") +
  # Adds title to graph, axis names to x and y axis and fills in plot with values in
  # Growth
  labs(fill = "Growth", title = "Population Growth in 2017", x="Longitude", y="Latitude")+
  # Sets scale limits
  xlim(-25,50) + ylim(35,70)

```



```
##                               sysname
##                               "Linux"
##                               release
##                               "5.4.0-66-generic"
##                               version
## "#74-Ubuntu SMP Wed Jan 27 22:54:38 UTC 2021"
##                               nodename
##                               "machinal28-linux64"
##                               machine
##                               "x86_64"
##                               login
##                               "machinadmz"
##                               user
##                               "machinadmz"
##                               effective_user
##                               "machinadmz"
```